

Multiagent Reinforcement Learning

Michael Krumbick, Geoffrey Liu

1 Main Findings So Far

1. AlphaRank is *slow*.
2. Noisy evaluation matrices degrade performance, but still work.
3. Good meta-solvers induce distributions that are very different from the baseline strategies (Uniform, biased Uniform, etc).

2 Background and Related Works

2.1 MDPs

A **Markov decision process (MDP)** [1] is a common formulation of decision making problem. Each MDP can be defined by a four-item tuple. $\langle S, A, P, R \rangle$

1. S : The space of valid states.
2. A : The space of valid actions.
3. $P_a(s, s')$: A transition function.
4. $R_a(s, s')$: A reward function.

The solution to the MDP will come in the form of a policy π which represents a mapping $S \rightarrow \Delta(A_s)$, or from the state space to a distribution over action space. The policy deemed optimal will obtain the highest expected reward, with some discount.

$$\pi^* = \arg \max_{\pi \in \Pi} \sum_{s \in S} d^\pi(s) \mathbb{E}_{a \sim \pi(s)} r(s, a) \quad (1)$$

Reinforcement learning is one method of finding these policies that has seen success. This family of methods has seen increased attention Algorithms such as ...

A **normal form game** can be defined (N, A, u)

1. N , the number of players
2. A , the number of players
3. u , the number of players

Algorithm 1 Double Oracle

```
actions  $\leftarrow$  RandomAction()  
 $M \leftarrow$  ComputePayoff(actions)  
for i = 1:N do  
     $\pi_i \leftarrow$  MetaSolver( $M$ )  
    newAction  $\leftarrow$  ResponseOracle( $\pi_i$ , actions)  
     $M_i \leftarrow$  UpdatePayoff(actions, newAction)  $\triangleright$  Fill in the new rows and columns for the new agent  
    actions  $\leftarrow$  append(actions, newAction)  
end for
```

Every time you add a new agent, you need to fill in $2n + 1$ cells.

Kuhn Poker is a simplified version of poker. The deck consists of three cards, the King, Queen and Jack,

Algorithm 2 PSRO

```
agents  $\leftarrow$  RandomAgents()  
M  $\leftarrow$  ComputePayoff(agents)  
for i = 1:N do  
   $\pi_i \leftarrow$  MetaSolver(M)  
  newAgent  $\leftarrow$  ResponseOracle( $\pi_i$ , agents)  
   $M_i \leftarrow$  UpdatePayoff(agents, newAgent)  $\triangleright$  Fill in the new rows and columns for the new agent  
  agents  $\leftarrow$  append(agents, newAgent)  
end for
```

Name	Time Complexity	Convergence
Uniform	$O(1)$	-
Regret Minimization	$O(n)$	-
Nash	$O(n^p)$	-
AlphaRank	$O(n^p)$	-

Table 1: Meta Solvers for Zero Sum Games

References

- [1] Martin L. Puterman. Markov decision processes: Discrete stochastic dynamic programming. In *Wiley Series in Probability and Statistics*, 1994.
- [2] Eilon Solan and Nicolas Vieille. Stochastic games. *Proceedings of the National Academy of Sciences*, 112(45):13743–13746, 2015.